

Computer Vision

Paulo Dias, António Neves






- Edges
 - Introduction
 - Edge detection
- Lines and corners
 - Line detection operators
 - Hough Transform
 - Harris corner detector
 - Other feature detectors




- Edges
 - Introduction
 - Edge detection
- Lines and corners
 - Line detection operators
 - Hough Transform
 - Harris corner detector
 - Other feature detectors



- Edges are useful to capture important events and changes in properties of the images/world
- Edge detection is difficult
 - noise
 - non ideal edges




Ideal edge




non ideal edge

- Edges correspond to
 - discontinuities in depth,
 - discontinuities in surface orientation,
 - changes in material properties,
 - variations in scene illumination.
- How to detect the relevant edges?



Edge Detection - example




Burger and Burge




- Edges
 - Introduction
 - **Edge detection**
- Lines and corners
 - Line detection operators
 - Hough Transform
 - Harris corner detector
 - Other feature detectors

- Typical 2 steps for edge detection:
 - Apply a **mask** (to approximate a derivative)
 - Aggregate detected pixels (**edgels**) in edges



- Derivatives are used to detect edges

- 1st derivative
 - $>$ or $<$ 0 depending on $I(x)$ variation
 - $=0$ in areas of same intensity
- 2nd derivative
 - $=0$ in both positive and negative edges






- 1st derivative



(a)





$$f'(x) = \frac{df}{dx}(x)$$

Edge detection – 1st derivative



- 1st derivative – simple approximation



$$\frac{df}{du}(u) \approx \frac{f(u+1) - f(u-1)}{2} = 0.5 \cdot (f(u+1) - f(u-1))$$



- Partial derivatives and gradient

$$H_x^D = \begin{bmatrix} -0.5 & \mathbf{0} & 0.5 \end{bmatrix} = 0.5 \cdot \begin{bmatrix} -1 & \mathbf{0} & 1 \end{bmatrix}$$

$$H_y^D = \begin{bmatrix} -0.5 \\ \mathbf{0} \\ 0.5 \end{bmatrix} = 0.5 \cdot \begin{bmatrix} -1 \\ \mathbf{0} \\ 1 \end{bmatrix}$$


$$\nabla I(u, v) = \begin{bmatrix} \frac{\partial I}{\partial u}(u, v) \\ \frac{\partial I}{\partial v}(u, v) \end{bmatrix}$$

$$|\nabla I|(u, v) = \sqrt{\left(\frac{\partial I}{\partial u}(u, v)\right)^2 + \left(\frac{\partial I}{\partial v}(u, v)\right)^2}$$


- Partial derivatives and gradient



Original image



H_x^D






H_y^D



$|\nabla I|(u, v)$

- Derivatives Operators
 - Image **gradient** points into the **direction** of larger intensity variation

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

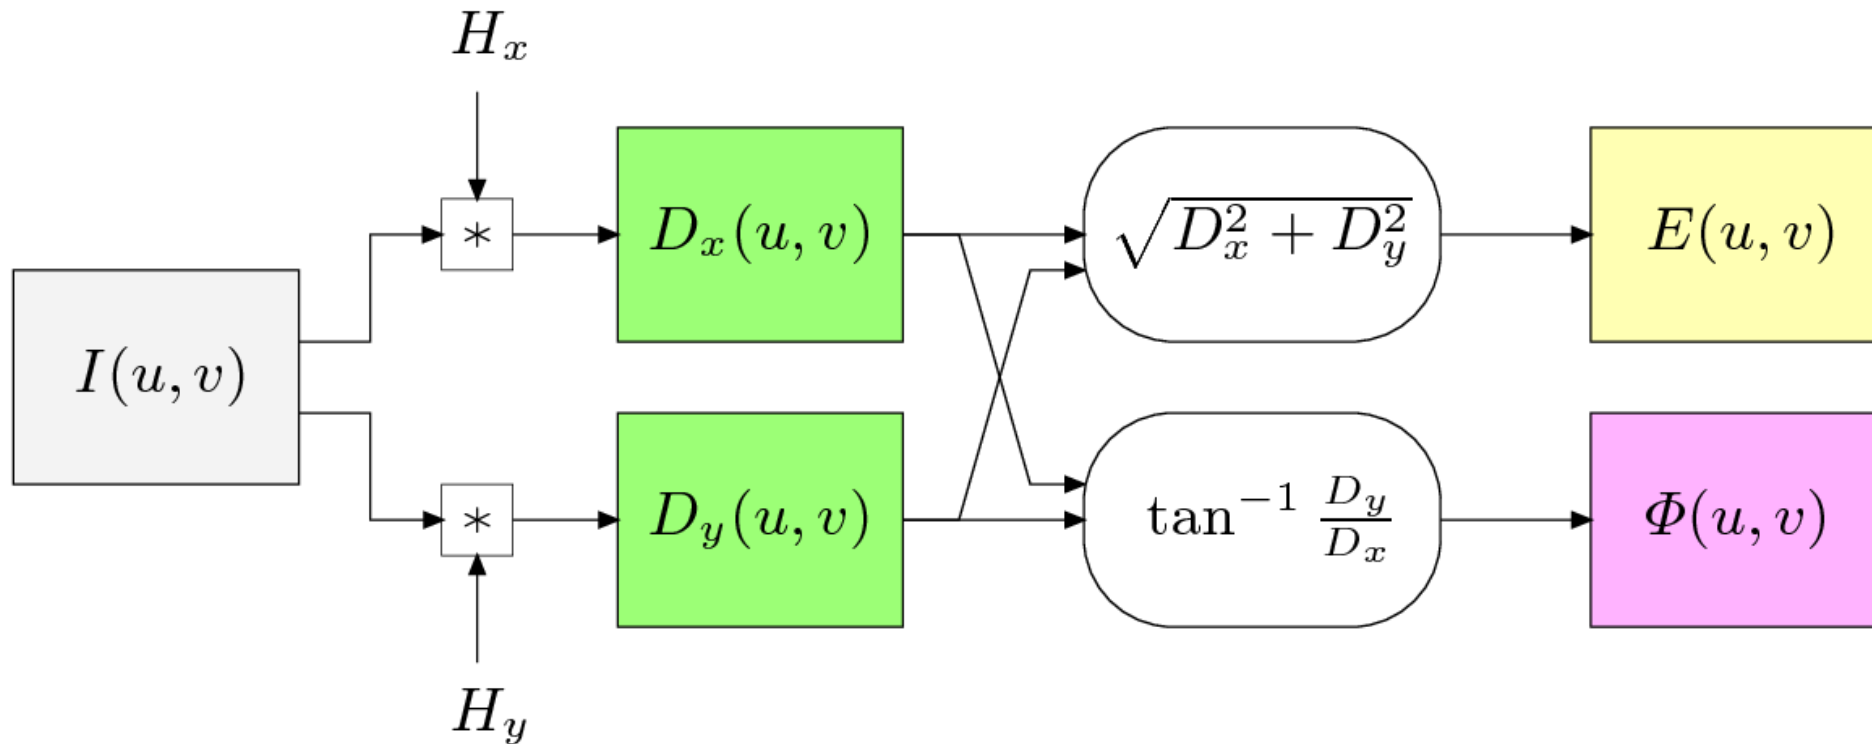
Gradient amplitude

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$


Gradient direction



- How to use operators?




- How to use operators?



- Roberts operator
 - Simple, fast but very noise sensitive



$$D_1 = I * H_1^R$$



$$D_2 = I * H_2^R$$

$$H_1^R = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} \quad \text{and} \quad H_2^R = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$

Burger and Burge

- Prewitt operator



$$H_x^P = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad \text{and} \quad H_y^P = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

-1	-1	-1
0	0	0
-1	-1	-1

Horizontal

-1	0	-1
-1	0	-1
-1	0	-1

Vertical



- Sobel operator



$$H_x^S = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad \text{and} \quad H_y^S = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

-1	-2	-1
0	0	0
1	2	1

-1	0	1
-2	0	2
-1	0	1

Horizontal

Vertical






- Sobel operator



Original Image


X – Direction Kernel

-1	0	1
-2	0	2
-1	0	1



Y – Direction Kernel

-1	-2	-1
0	0	0
1	2	1



Resulting image



- Compass Edge Detection
 - alternative to gradient edge detection (Roberts and Sobel operators).
- Usually outputs two images
 - Gradient magnitude
 - edge orientation
- Gradient is estimated in eight (for a 3 x 3 convolution mask) possible orientation (from 0° [vertical] to 315° in steps of 45°).
- The convolution result of greatest magnitude indicates the gradient direction



- Extended-Sobel Operator

$$\begin{aligned} H_0^{\text{ES}} &= \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, & H_1^{\text{ES}} &= \begin{bmatrix} -2 & -1 & 0 \\ -1 & 0 & 1 \\ 0 & 1 & 2 \end{bmatrix}, \\ H_2^{\text{ES}} &= \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}, & H_3^{\text{ES}} &= \begin{bmatrix} 0 & -1 & -2 \\ 1 & 0 & -1 \\ 2 & 1 & 0 \end{bmatrix}, \\ H_4^{\text{ES}} &= \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}, & H_5^{\text{ES}} &= \begin{bmatrix} 2 & 1 & 0 \\ 1 & 0 & -1 \\ 0 & -1 & -2 \end{bmatrix}, \\ H_6^{\text{ES}} &= \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}, & H_7^{\text{ES}} &= \begin{bmatrix} 0 & 1 & 2 \\ -1 & 0 & 1 \\ -2 & -1 & 0 \end{bmatrix}. \end{aligned}$$



- Laplacian operator
 - Second derivative approximation of ∇^2

4-neighborhood

$$h = \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

8-neighborhood

$$h = \begin{bmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Edge detection with **first derivative** are **noise** sensitive and object dependent
 - The first derivative of the image function should have an **extremum** at the position corresponding to the edge
 - It is much **easier and more precise** to find a **zero-crossing** position than an extremum.



- Canny objectives
 - Good location (zero crossing)
 - Minimize weak edges




- Canny Edge Detector (1986)
 - Process in five steps:
 1. Gaussian filter to smooth and remove noise
 2. Find intensity gradients of the image (Sobel operator)
 3. Non-maximum suppression
 4. Double threshold to determine potential edges



High threshold for strong pixels

Low threshold for non-relevant pixels
 5. Edge Tracking by Hysteresis


Transform weak into strong pixels, if at least one neighboring pixels is processed as strong






- Canny Edge Detector (1986)
 - Process in five steps:
 1. Gaussian filter to smooth and remove noise



- Canny Edge Detector (1986)
 - Process in five steps:
 1. Convert image to grayscale
 2. Find intensity gradients of the image (Sobel operator)





- Canny Edge Detector (1986)
 - Process in five steps:
 3. Non-maximum suppression



- Canny Edge Detector (1986)
 - Process in five steps:
 4. Double threshold to determine potential edges


High threshold for strong pixels

Low threshold for non-relevant pixels




- Canny Edge Detector (1986)
 - Process in five steps:
 5. Edge Tracking by Hysteresis

Transform weak into strong pixels, if at least one neighboring pixels is processed as strong



Images from <https://towardsdatascience.com/canny-edge-detection-step-by-step-in-python-computer-vision-b49c3a2d8123>

Edge detection - Canny




Original Image



1D convolution (x,y) with Gaussian smoothing $\sigma = 6$



Canny edges



Edge detection - Canny



Original



$\sigma = 1.0$



$\sigma = 2.0$



$\sigma = 4.0$



$\sigma = 8.0$



$\sigma = 16.0$

Burger and Burge

Edge detection - comparison




- Possible criteria:
 - Number of weak/false edges
 - Connectivity
 - ...




Original



Roberts




Prewitt



Sobel



Laplacian of Gaussian



Canny ($\sigma = 1.0$)



- Edges
 - Introduction
 - Edge detection
- Lines and corners
 - Line detection operators
 - Hough Transform
 - Harris corner detector
 - Other feature detectors



- Same rationale of detecting "roof" like profiles along "strategic" orientations: 0° ; 45° ; 90° ; 135° [see compass]
- Convolution Kernels

$$h_1 = \begin{bmatrix} -1 & -1 & -1 \\ 2 & 2 & 2 \\ -1 & -1 & -1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 2 & -1 & -1 \\ -1 & 2 & -1 \\ -1 & -1 & 2 \end{bmatrix} \quad h_3 = \begin{bmatrix} -1 & 2 & -1 \\ -1 & 2 & -1 \\ -1 & 2 & -1 \end{bmatrix}$$

- Lines detected this way are collections of edges. Most of the time non single pixel wide edges.
- Necessary to introduce line thinning algorithms




- Edges
 - Introduction
 - Edge detection
- Lines and corners
 - Line detection operators
 - **Hough Transform**
 - Harris corner detector
 - Other feature detectors


Hough transform (1962)




- Technique for having edges “vote” for plausible line locations
- Represent line edges in polar coordinates (r, θ) in the Hough space



Line edge in original image




Conversion to (r, θ) representation




Line representation in Hough space


- Line representation in Hough space



θ	r
15	189.0
30	282.0
45	355.7
60	407.3
75	429.4




θ	r
15	318.5
30	376.8
45	407.3
60	409.8
75	385.3




θ	r
15	419.0
30	443.6
45	438.4
60	402.9
75	340.1

- Higher cell values in Hough accumulator are the Hough parameters of the lines for which angle and distance can be determined



Input Image



Rendering of Transform Results




- Classical Hough transform for line identification
- Extended to identifying positions of other shapes as circles or ellipses.





- Edges
 - Introduction
 - Edge detection
- Lines and corners
 - Line detection operators
 - Hough Transform
 - **Harris corner detector**
 - Other feature detectors

- Corners in images can be located using local detectors;
 - Input to the corner detector is the Gray-level image
 - Output is the image in which values are proportional to the likelihood that the pixel is a corner.
 - Interest points are obtained by thresholding the result of the corner detector.
- Edge detectors themselves are not stable at corners.
 - Gradient at the tip is ambiguous





- One of the earliest corner detection - 1979
- Corner: point with low self-similarity
 - Tests pixels as corners considering similarity between nearby, largely overlapping patches.
 - Similarity is measured by taking the sum of squared differences (SSD) between the corresponding pixels of two patches



- Auto-correlation based
- Improvement upon Moravec's corner detector
- Use a sliding window W patch and estimate the sum of square differences of the discriminant function:

$$N = \begin{bmatrix} \sum_{\text{window}} f_r^2(r, c) & \sum_{\text{window}} f_r(r, c) \cdot f_c(r, c) \\ \sum_{\text{window}} f_r(r, c) \cdot f_c(r, c) & \sum_{\text{window}} f_c^2(r, c) \end{bmatrix} \quad \begin{array}{l} f_r(r, c) : \text{horizontal gradient} \\ f_c(r, c) : \text{vertical gradient} \end{array}$$

- Compute smallest eigenvalue of the structure tensor:

$$\lambda_{\min} \approx \frac{\lambda_1 \lambda_2}{(\lambda_1 + \lambda_2)} = \frac{\det(M)}{\text{tr}(M)}$$

with the trace $\text{tr}(M) = m_{11} + m_{22}$.



- Other corner detector:
 - Kitchen - 82
 - Harris - 88
 - Deriche - 90
 - Mehrotra - 90
 - Schmid - 98
 - Smith – 98
 - ...




- Edges
 - Introduction
 - Edge detection
- Lines and corners
 - Line detection operators
 - Hough Transform
 - Harris corner detector
 - Other feature detectors



- Several feature descriptors more or less invariant to scale, rotation, affine transformations:
 - Histogram based (use histogram of oriented gradient)
 - SIFT - Scale invariant feature transform
 - SURF - Speeded-Up Robust Features
 - GLOH - Gradient Location and Orientation Histogram
 - HOG - Histogram of Oriented Gradients
 - Compact descriptors (use binary strings comparing pairs of intensity images)
 - BRIEF - Binary Robust Independent Elementary Features
 - FAST - Features from accelerated segment test
 - ORB - Oriented FAST and Rotated BRIEF
 - BRISK - Binary Robust invariant scalable keypoints

- SIFT Basic idea:
 - Take 16x16 square window around detected feature
 - Compute edge orientation (angle of the gradient) for each pixel
 - Throw out weak edges (threshold gradient magnitude)
 - Create histogram of surviving edge Segment Test



- SIFT Full version:
 - Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
 - Compute an orientation histogram for each cell
 - 16 cells * 8 orientations = 128 dimensional descriptor

